

3 Fine-tuning your model

0%

Having trained your model, your next task is to evaluate its performance. What metrics can you use to gauge how good your model is? So far, you have used accuracy for classification and R-squared for regression. In this chapter, you will learn about some of the other metrics available in scikit-learn that will allow you to assess your model's performance in a more nuanced manner. You will then learn to optimize both your classification as well as regression models using hyperparameter tuning.

▶	How good is your model?	50 xp
</>	Metrics for classification	100 xp
▶	Logistic regression and the ROC curve	50 xp
</>	Building a logistic regression model	100 xp
</>	Plotting an ROC curve	100 xp
≡	Precision-recall Curve	50 xp
▶	Area under the ROC curve	50 xp
</>	AUC computation	100 xp
▶	Hyperparameter tuning	50 xp
</>	Hyperparameter tuning with GridSearchCV	100 xp
</>	Hyperparameter tuning with RandomizedSearchCV	100 xp
▶	Hold-out set for final evaluation	50 xp
≡	Hold-out set reasoning	50 xp
</>	Hold-out set in practice I: Classification	100 xp
</>	Hold-out set in practice II: Regression	100 xp

HIDE CHAPTER DETAILS